Ve 280

Programming and Introductory Data Structures

Interfaces; Invariants

Learning Objectives:

Understand what interfaces are and how to implement them in C++

Understand better what invariants are and how to use them to prevent some bugs

Outline

- Interfaces

- Invariants

ADTs

Recall

- Recall the two main advantages of an ADT:

  1. <u>Information hiding</u>: we don't need to know the details of **how** the object is represented, nor do we need to know how the operations on those objects are implemented.

  2. <u>Encapsulation</u>: the objects and their operations are defined in the same place; the ADT combines both data and operation in one entity.

ADTs

Recall

- To the caller, an ADT is only an **interface.**

  - **Interface**: the contract for using things of this type.

- Once you have an interface, you can pick from among many possible implementations as long as you satisfy the contract.

```
class IntSet { // a mutable set of integers
  public:
    void insert(int v); // this + {v}
    void remove(int v); // this – {v}
    bool query(int v); // does v exist in this?
    int  size(); // return |this|
};
```

Interfaces

Separating out the details

- The class mechanism, as we've used it so far, has one shortcoming:
  - It mixes details of the **implementation** with the definition of the **interface**.

Interfaces

Separating out the details

- Recall that the implementation of a class includes:
    1. Data members
    2. Method implementations
- The method implementations can be written separately from the class definition and are usually in two separate files.
    - Class definition in .h file; method implementation in .cpp file.
- Unfortunately, the **data members** still must be part of the class definition (in .h file).
    - Since any programmer using an `IntSet` must see that definition, those programmers know something about the implementation.

Interfaces

Separating out the details

- Having data objects in the definition has two undesirable effects:
    1. It complicates the class definition, making it harder to read and understand.
    2. It communicates information to the programmer that s/he shouldn't know.

- The second problem can have very drastic consequences.
    - If a programmer using your class (mistakenly) makes an assumption about a "guarantee" that your implementation provides, but the interface doesn't promise, he is in trouble when you change the implementation.

Interfaces

Separating out the details

- **<u>Question</u>**: How can you provide a class definition that carries no implementation details (i.e., data members) to the client programmer, yet still has interface information?

- **<u>Answer</u>**: Create an "interface-only" class as a **base class**, from which an implementation can be **derived**.
  - <u>Note</u>: classes **must** contain their data members, so this class **cannot** have a real implementation!
  - Such a base class is called an **Abstract Base Class**, or sometimes a **Virtual Base Class**, because we're going to leverage virtual methods to do it.

Interfaces

Creating an abstract base class

- To create an abstract base class, we first provide an "interface-only" definition of `IntSet`.

- Because there will be no implementation, we need to declare its methods in a special way:
  - Declare each method as a virtual function.
  - "Assign" a zero to each of these virtual functions.

```cpp
class IntSetFull { };

class IntSet {
  // OVERVIEW: mutable set of ints with bounded size
public:
  virtual void insert(int v) = 0;
    // MODIFIES: this
    // EFFECTS: set=set+{v}, throws IntSetFull if full
  virtual void remove(int v) = 0;
    // MODIFIES: this
    // EFFECTS: set=set-{v}
  virtual bool query(int v) = 0;
    // EFFECTS: returns true if v is in set,
    //          false otherwise
  virtual int  size() = 0;
    // EFFECTS: returns |set|
};
```

Interfaces

Creating an abstract base class

```cpp
class IntSetFull { };
class IntSet {
  public:
  virtual void insert(int v) = 0;
  virtual void remove(int v) = 0;
  virtual bool query(int v) = 0;
  virtual int  size() = 0;
};
```

- These functions are called **pure virtual functions** and are declared not to exist.
- Think about them as a set of **function pointers**, all of which point to NULL.

Interfaces

Creating an abstract base class

```cpp
class IntSetFull {};
class IntSet {
  public:
  virtual void insert(int v) = 0;
    // MODIFIES: this
    // EFFECTS: set=set+{v}, throws
    //          IntSetFull if full
  virtual void remove(int v) = 0;
  virtual bool query(int v) = 0;
  virtual int  size() = 0;
};
```

Note the use of `IntSetFull` as an "exception type". It is used as something convenient to `throw` instead of some random `int`.

Interfaces

Abstract base classes

- A class with one or more Pure Virtual Functions is an **abstract** class.
- You **cannot** create **any instances** of an abstract class, because there are no implementation.

- For example, the following fails:
```
IntSet s;
```

- However, you can always define **references** and **pointers** to an abstract class, so these are both legal:
```
IntSet &r = <something>;
IntSet *p;
```

Interfaces

Abstract base classes

- Abstract base classes aren't very interesting without some derivative of `IntSet` to actually provide an implementation.
- This is done with a simple derived class:

```
const int MAXELTS = 100;
class IntSetImpl : public IntSet {
    int elts[MAXELTS];
    int numElts;
public:
    IntSetImpl();
    void insert(int v);
    void remove(int v);
    bool query(int v);
    int  size();
};
```

Note: The implementation has data members.

In general, besides new function members, a derived class can also have new data members.

Interfaces

Abstract base classes

- Abstract base classes aren't very interesting without some derivative of `IntSet` to actually provide an implementation.
- This is done with a simple derived class:

```
const int MAXELTS = 100;
class IntSetImpl : public IntSet {
  int elts[MAXELTS];
  int numElts;
public:
  IntSetImpl();
  void insert(int v);
  void remove(int v);
  bool query(int v);
  int  size();
};
```

Note: This implementation could be **either** the sorted or unsorted versions.

Interfaces

Abstract base classes

- Abstract base classes aren't very interesting without some derivative of `IntSet` to actually provide an implementation.
- This is done with a simple derived class:

```
const int MAXELTS = 100;
class IntSetImpl : public IntSet {
   int elts[MAXELTS];
   int numElts;
public:
   IntSetImpl();
   void insert(int v);
   void remove(int v);
   bool query(int v);
   int  size();
};
```

Note: the derived class has to implement the constructor. In the past, it was always in the base class.

It can't be there, because the base class has no implementation to construct!

In principle, should the implementation code of the derived class of an abstract class be provided to its user?

Select all the correct answers.

- **A.** Yes, so the user understands how the abstract class is implemented.

- **B.** Yes, so the constructor can be called by the user.

- **C.** No, it would go against the spirit of an ADT.

- **D.** No, no file related to the implementation is needed by the user.

Interfaces

Abstract base classes

- The interface (the abstract base class) is typically defined in a public header (*.h) file
  - Users of the **interface** include the *.h file.
- The implementation (**the derived class**) is defined in a source (*.cpp) file
  - Users of the interface only *link* against (i.e., compile the file into object code and link with other object codes)
- So, a user of the `IntSet` abstraction **never sees** the definition for class `IntSetImpl`.
- The only thing that remains is to give users the means to create a new `IntSet`:
  - However, they can't do it in the normal way: `IntSet s;`
  - Also, they can't create objects of the derived class, because its definition is **not visible** to them.

Interfaces

Abstract base classes

- If only one instance of the class is needed, the *.h file typically includes the following prototype for an access function:

```
// header file
IntSet *getIntSet();
  // EFFECTS: returns a pointer
  //          to the IntSet
```

- The *.cpp file defines a single, static instance (only visible to the *.cpp file) of the implementation and body of the access function:

```
// source file
static IntSetImpl impl;
IntSet *getIntSet() {
  return & impl;
}
```

Interfaces

Abstract base classes

- If only one instance of the class is needed, the *.h file typically includes the following prototype for an access function:

```
// header file
IntSet *getIntSet();
  // EFFECTS: returns a pointer
  //          to the IntSet
```

- The *.cpp file defines a single, static instance (only visible to the *.cpp file) of the implementation and body of the access function:

```
// source file
static IntSetImpl impl
IntSet *getIntSet() {
   return & impl;
}
```

Note: Now the user can do the following and it will be valid:

```
IntSet *s = getIntSet();
```

Interfaces

Abstract base classes

- If more than one instance of the class is needed, we need to provide a function that creates them **dynamically**…
  - You will see how to do this later.

Outline

- Interfaces

- Invariants

Invariants

- An invariant is a set of conditions that must always evaluate to true at certain well-defined points; otherwise, the program is incorrect.

- For ADT, there is so called **representation invariant**.

Invariants

- A **<u>representation invariant</u>** applies to the data members of ADT.

- It describes the conditions that must hold on those members for the representation to correctly implement the abstraction.

- It must hold <u>immediately before exiting each method</u> of that implementation – including the constructor.
  - Example: insert() member of IntSet.
  - This is called **establishing the invariant**.

Invariants

Representation Invariants

- Each method in the class can assume that the invariant is true **on entry** <u>if</u>:
  - The representation invariant holds <u>immediately before exiting each method</u> (including the constructor), **and**
  - Each data element is truly private.

- This is true because the only code that can change the data members belongs to the methods of that class, and those methods always establish the invariant.

Invariants

Representation Invariants

- We've seen two examples of representation invariants, both applied to the private data members of an `IntSet` representation:

```
int elts[MAXELTS];
int numElts;
```

- For the unsorted version, the invariant is:
  - The first `numElts` members of `elts` contain the integers comprising the set, with no duplicates.
- For the sorted version, the invariant is:
  - The first `numElts` members of `elts` contain the integers comprising the set, from lowest to highest, with no duplicates.

Invariants

Representation Invariants

- We used these invariants to write the methods of each implementation.
- For example:

```
insert(int v)         // unsorted version
  if v not in elts // don't allow duplicates
    elts[numElts] = v  // this breaks invariant
    numElts++          // this restores it


insert(int v)         // sorted version
  if v not in elts // don't allow duplicates
    make gap in array // this breaks invariant
    elts[gap] = v  // restore elts invariant
    numElts++      // restore numElts invariant
```

Invariants

Representation Invariants

- The representation invariant plays a crucial role in implementing an abstract data type.

- Before writing a **single** line of code, write down the rep invariant!

- That tells you **how** to write each method.

- Essentially, for each method, you should:
  - Do the work of the method (i.e. insert)
  - Repair the invariants you broke

Invariants

Checking for Representation Invariants

- Invariants can also be coded, to check the sanity of the structure.

- For even moderately complicated data structures, it is worth writing a function to check for invariants.

- In the IntSet case, we **can** check to see if the array satisfies the respective invariants such as there is no duplication or the array is sorted.

Invariants

Checking for Representation Invariants

- Use sorted representation for example. We will write the following function to check the invariants:

```
bool strictSorted(int a[], int size)
  // REQUIRES: a has size elements
  // EFFECTS: returns true if a is sorted
  //               with no duplicates
```

- How can we tell if an array is sorted with no duplicates?

  - If size ≤ 1, the array is sorted with no duplicates.

  - If size > 1, then the array must satisfy
    a[0] < a[1] < … < a[size-1]

Invariants

Checking for Representation Invariants

```
bool strictSorted(int a[], int size)  {
  // REQUIRES: a has size elements
  // EFFECTS: returns true if a is sorted
  //            with no duplicates

  if (size <= 1) return true;

  for (i=0; i<size-1; i++){
    if (a[i] >= a[i+1]) {
      return false;
    }
  }
  return true;
}
```

Invariants

Checking for Representation Invariants

- Writing these "checker" functions is very useful – you can use them for **defensive programming**.

- So, you can write a **<u>private</u>** method to check whether all invariants are true (**before exiting**, or after entering, each method):

```
bool repOK();
// EFFECTS: returns true if the
//          rep. invariants hold
```

- For the sorted version, `repOK` would be:

```
bool repOK() {
    return strictSorted(elts, numElts);
}
```

Invariants

Checking for Representation Invariants

- Next, add the following code **right before** returning from any function that modifies any of the representation:

```
assert(repOK());
```

- If you are truly paranoid, you can write the same line at the **beginning** of every method, too; this checks that the assumption the method relies on is true.

References

- **Problem Solving with C++ (8th Edition)**, by *Walter Savitch*, Addison Wesley Publishing (2011)
  - Chapter 10.4 Introduction to Inheritance
  - Chapter 15.1 Inheritance Basics
  - Chapter 15.3 Virtual Functions in C++